

Orchestrating Microbiome Analysis with Bioconductor

Tuomas Borman^{1*}, Giulio Benedetti¹, Geraldson Muluh¹,
Aura Raulo^{3,4}, Benjamin Valderrama^{5,6}, Artur Sannikov²,
Stefanie Peschel⁷, Yihan Liu^{8,9,10,11,12}, Rasmus Hindström¹,
OMA consortium¹⁹, Katariina Pärnänen¹³,
Christian L. Müller^{7,14,15}, Aki S. Havulinna^{1,16},
Sudarshan Shetty¹⁷, Marcel Ramos¹⁸, Domenick J. Braccia¹⁹,
Héctor Corrada Bravo¹⁹, Felix G. M. Ernst¹⁹, Levi Waldron¹⁸,
Thomaz F. S. Bastiaanssen^{20,21†}, Himel Mallick^{8,22,23†}, Leo Lahti^{1*}

¹Department of Computing, University of Turku, Turku, Finland.

²Department of Life Technologies, University of Turku, Turku, Finland.

³Department of Biology, University of Oxford, Oxford, United Kingdom.

⁴Queen's College, University of Oxford, Oxford, United Kingdom.

⁵Department of Anatomy and Neuroscience, University College Cork,
Cork, Ireland.

⁶APC Microbiome Ireland, University College Cork, Cork, Ireland.

⁷Department of Statistics, LMU München, Munich, Germany.

⁸Department of Population Health Sciences, Weill Cornell Medicine,
Cornell University, New York, USA.

⁹Department of Biomedical Informatics & Data Science, Cornell
University, New York, USA.

¹⁰Department of Biostatistics, Cornell University, New York, USA.

¹¹Department of Biomedical Informatics & Data Science, Yale University,
New Haven, USA.

¹²Department of Biostatistics, Yale University, New Haven, USA.

¹³Department of Microbiology, University of Helsinki, Helsinki, Finland.

¹⁴, Helmholtz Munich, Munich, Germany.

¹⁵, Flatiron Institute, New York, USA.

¹⁶, Institute for Molecular Medicine Finland, FIMM-HiLIFE, Helsinki, Finland.

¹⁷, Danone Research and Innovation, Utrecht, Netherlands.

¹⁸Department of Epidemiology and Biostatistics, CUNY Graduate School of Public Health and Health Policy, New York, USA.

¹⁹, No affiliation.

²⁰Department of Psychiatry, Amsterdam University Medical Centres, Amsterdam, Netherlands.

²¹Amsterdam Public Health, Amsterdam University Medical Centres, Amsterdam, Netherlands.

²²Division of Gastroenterology and Hepatology, Weill Cornell Medicine, Cornell University, New York, USA.

²³Department of Statistics and Data Science, Cornell University, New York, USA.

*Corresponding author(s). E-mail(s): tuomas.v.borman@utu.fi;
leo.lahti@utu.fi;

[†]These authors contributed equally to this work.

Abstract

The expansion of microbiome research has led to the accumulation of interlinked datasets encompassing versatile taxonomic and functional assays. The analysis of increasingly large and heterogeneous multi-modal microbiome data would benefit from unified approaches supporting the design of modular data science workflows through interoperable methods. The Bioconductor project recently developed an optimized statistical programming framework for multi-assay data integration. Building on this foundation, we introduce a community-developed, open-source framework for microbiome data science. In contrast to the previous alternatives, the methodology is specifically designed to support joint analysis of hierarchical, interlinked, and heterogeneous multi-table datasets that are increasingly common in modern microbiome research. This framework encompasses open data, methods, tutorials, and an active online community, thereby

supporting standardized and reproducible data wrangling, joint analysis, and reporting. We have detailed the functionality and usage in the OMA book (<https://bioconductor.org/books/release/OMA>), which offers guidance for prospective users and contributors.

Keywords: microbiome, Bioconductor, data science, data integration

Introduction

Modern data science applications critically rely on open-source methods created by the research community^{1;2}. Open software ecosystems have become central innovation hubs, driving technological and cultural transformation³. The Bioconductor project is a significant distribution channel for open data science methods in computational life sciences. Its vast developer community maintains over 2,300 quality-controlled R packages for various fields of life science informatics⁴. Additionally, Bioconductor has evolved into a global community that supports data science collaboration and training⁵, playing an increasingly important role in advancing computational skills that are essential for modern microbiology education⁶.

Microbiome research focuses on how microbial communities vary, function and interact with their environment and hosts⁷. The technologies used in this field generate vast amounts of high-dimensional and hierarchically structured data, commonly in the form of microbial DNA sequences detected in a set of samples. Following the rapid expansion of microbiome research and technological advances, metagenomic sequencing is complemented by other types of high-throughput measurements. This has led to the accumulation of multi-modal datasets—collections that combine different types of molecular data. These interlinked datasets encompass diverse taxonomic and functional assays, ranging from microbial abundances and phylogenetic relationships to functional profiles. These are summarized into tabular data sets whose typical sizes range up to tens of thousands of selected features and samples for downstream statistical analyses.

Microbiome data scientists increasingly face the practical challenges of integrating heterogeneous and high-dimensional observations across multiple data modalities into scalable and reproducible workflows^{8–10}. The unique statistical properties and complexity of microbiome data necessitate specialized approaches^{7;11–14}, yet the diversity of proposed solutions and inconsistent outcomes^{15–17} can make it difficult to identify the optimal methods and build interoperable workflows^{18;19}. Accordingly, the field can benefit for standardization and evidence-based best practices⁷.

Open microbiome data science *frameworks* have emerged to provide sets of interoperable methodologies in specific computational environments^{20–23}, and a broader ecosystem of digital resources, engaged user communities and collaboration culture

that extend the capacities far beyond the underlying technical framework. However, frameworks that are focused solely on microbiome data science fall short on addressing the need to integrate data across multiple modalities, such as metagenomics, metabolomics, and host genomics. Bioconductor integrates the methods of microbiome data science into a wider statistical programming framework, where interoperable methods for single-cell omics²⁴, spatial transcriptomics²⁵, metabolomics and proteomics²⁶, and other fields are concurrently developed.

This approach can introduce some challenges, such as dependency management, package compatibility, and differences in package stability. These issues are addressed through coordinated mechanisms that help to ensure the robustness of the software ecosystem. First, submitted packages undergo peer review, ensuring relevance and adherence to standards such as unit tests, reproducible documentation, and working examples. Secondly, Bioconductor follows a synchronized, semiannual release cycle, during which packages are frozen and mutually compatible. Thirdly, all packages are continuously tested within Bioconductor’s automated build system, which ensures that breaking changes are rapidly detected before release. Each package has an active maintainer who is notified of failures and responsible for resolving them in a timely manner. Finally, community guidelines and continuous integration practices, together with shared data structures such as (Tree)SummarizedExperiment, further improve interoperability between packages.

Here, we introduce a stable and optimized data science ecosystem supporting the integration and analysis of multi-modal datasets in microbiome research. It provides harmonized data import and representation, and fast, robust methods for transforming abundance tables into biological insights. These methods have been extensively tested and distributed through mature R/Bioconductor packages (<https://bioconductor.org/packages>), supporting the community-driven development and integration within the broader Bioconductor ecosystem. To promote evidence-based best practices, we introduce the online book *Orchestrating Microbiome Analysis with Bioconductor* (OMA), which represents a comprehensive compilation of established approaches and emerging trends in microbiome data science using the R/Bioconductor ecosystem.

Data infrastructure

Data scientists routinely deal with interlinked assays, hierarchical data (*e.g.*, ontologies, trees), relational data (*e.g.*, networks), and supporting side information. This side information can include sample attributes such as treatment group, collection site, or time point. The ability to integrate heterogeneous data elements is essential in biological research, where the individual components of a system can seldom be studied in isolation. However, this growing complexity introduces additional overhead, as researchers must track samples and features across multiple tables and manage an increasing number of data elements and non-trivial links between them. Moreover, input data formats vary across methods, leading to technical complications and time lost on data wrangling. Bioconductor addresses these challenges through standardized

data structures that are extended by an ecosystem of quality-controlled, interoperable methods. Adopting the framework will thus allow the analyst to avoid unnecessary data wrangling and invest more time in core analytical tasks.

Bioconductor data structures support statistical analysis of complex data combinations²⁷ and are widely adopted in omics research (Figure 1). They have been implemented through specialized object-oriented classes that integrate multiple data elements into a single structure known as a *data container*. The SummarizedExperiment (SE) class²⁷ is the primary data container underlying the Bioconductor ecosystem (Figure 2a and Table 1). Its position among the top 1% most-downloaded packages in Bioconductor reflects its widespread recognition and increasing adoption among developers. It provides a standardized solution to store tabular data by linking numeric abundance matrices with rich side information on the features/rows (*e.g.*, taxonomy), and samples/columns (*e.g.*, origin, collection time, host health).

The vast adoption and interoperability of SE-based data structures enables the transfer of methods across different fields of omics research, thereby promoting the development of modular and scalable data science methods and yielding improved overall quality control and user support⁵. The SE framework has been extended to similar data integration tasks, notably in single-cell and spatial omics through the SingleCellExperiment (SCE)²⁴ and SpatialExperiment²⁸ classes.





Figure 1: Adoption of the SummarizedExperiment (SE) data container among Bioconductor developers. The growth in the number of Bioconductor packages supporting SE over time is shown, categorized by their respective fields. Packages are counted by the date they were added to Bioconductor. Some packages added the SE support after their original release.

Hierarchical data structures. Microbiome data scientists frequently need to integrate information on feature phylogenies and sample hierarchies to accurately reflect

fine-scale variability in microbiome data. While maintaining interoperability with the broader Bioconductor ecosystem, the `TreeSummarizedExperiment` (TreeSE) class²⁹ extends the SE and SCE data structures to hierarchical datasets by incorporating both feature and sample organizations as tree structures (Figure 2b and Table 1). Beyond the core components of microbiome datasets, including taxa abundances, sample metadata, phylogenetic trees, and reference sequences, TreeSE provides a flexible foundation for multi-modal data integration, supporting the incorporation of functional predictions, resistome, metabolome, and other molecular profiles, as well as alternative feature representations. It can also store results from common analytical procedures, such as statistical transformations or agglomeration. By inheriting full compatibility with SE and SCE, TreeSE enables the development of modular, interoperable workflows across the Bioconductor ecosystem.


a

SummarizedExperiment (SE)
generic container linking multiple data tables (assays)



b

TreeSummarizedExperiment (TreeSE)
extension of SE incorporating hierarchical data structures



c

MultiAssayExperiment (MAE)
binding of omic-specific containers across multi-omics experiments




Figure 2: Integrative Bioconductor containers for microbiome data. (a) **SummarizedExperiment (SE)**²⁷ is a generic container for single-omics experiments. It accommodates *assays* that describe the feature abundances in each sample (*e.g.*, taxonomic units, resistance genes, or predicted functions), and tabular side information on each sample (*colData*) and feature (*rowData*), *e.g.*, the taxonomic classification. (b) **TreeSummarizedExperiment (TreeSE)**²⁹ supports hierarchical, multi-table microbiome analysis. In addition to the SE elements, TreeSE incorporates tree information on feature hierarchies, such as phylogenetic relations between the taxonomic units (*rowTree*) or host species (*colTree*), and results of dimensionality reduction (*reducedDims*). TreeSE can thus link multiple abundance tables and their transformed versions (*e.g.*, statistical transformation, taxonomic agglomeration, or dimensionality reduction) with tabular and hierarchical side information on the features (rows) and samples (columns). Additional slots for reference sequences and experiment metadata are available. (c) **MultiAssayExperiment (MAE)**³⁰ facilitates the integration of interlinked datasets representing different omics modalities linked by potentially non-trivial sample mappings. Figure adapted from³⁰.

Multi-assay data structures. Contemporary microbiome research frequently involves data across multiple measurement modalities for enhanced functional and mechanistic insights. This brings the need to find an appropriate representation for each data type and map the links between them. Whereas the TreeSE data container is broadly applicable to different modalities such as taxonomic, resistome, transcriptome, proteome, and metabolome profiling data, dedicated data containers have been developed for certain modalities by the Bioconductor community to address their specific characteristics; examples include single-cell sequencing²⁴ and host genomics³¹. Linking data across modalities presents an additional technical challenge, as multi-omics datasets are often sparse so that only a subset of the samples may be shared across multiple modalities³². In other cases, a single sample in one modality may link to several samples in another modality³³ (*e.g.*, due to technical or spatial replication). These scenarios require the ability to integrate multiple data modalities into a single object and create flexible sample mappings across them. The MultiAssayExperiment (MAE) data container³⁰ addresses these challenges by introducing a formal mapping layer (*sampleMap*) that links biological samples to their corresponding measurements in each modality (Figure 2c and Table 1). This design supports one-to-one, one-to-many, and partially overlapping sample relationships across experiments. Importantly, this mapping is also accessible for downstream operations. For instance, when subsetting a MAE object (*e.g.*, selecting samples with both microbiome and metabolome data), the framework automatically resolves the intersection of available samples and ensures consistent alignment across all modalities. Such structured handling of sample relationships is particularly important for integrative methods, as it reduces the risk of sample mismatches, a common source of irreproducibility in ad hoc multi-table analyses. To summarize, flexible mapping of samples across several data objects is a key for efficient data integration, method sharing, and interoperability.

Scalability. Enhanced scalability is one of the advantages of our framework, and essential with increasingly large and multi-modal data sets. This can be achieved by parallelization, out-of-memory analyses leveraging the built-in support for sparse and

delayed matrices, more efficient implementations through code optimization, executing costly computations in C++, and methodological advancements. In many cases, the support can be readily inherited from the underlying Bioconductor infrastructure, rather than newly implemented³⁴. For instance, the (Tree)SE and MAE containers are agnostic to assay representation and natively support sparse (dgCMatrix, SparseArray) and file-backed (HDF5Array, DelayedArray) assays²⁷. Many methods readily support parallelization; BiocParallel can support operations that decompose into independent blocks by samples, features, or other data elements. Non-decomposable operations, such as distance matrices and ordinations, must be realized in memory or parallelized by a dedicated algorithm: accelerated versions of common microbiome methods, including Faith phylogenetic alpha diversity³⁵ and striped UniFrac beta diversity³⁶, have been incorporated in the (Tree)SE-based framework through the mia package.

Table 1: Key elements of the Bioconductor data containers for microbiome analysis.

Slot	Content	SE	TreeSE	MAE ^a
Tables				
assays	tables of features (rows) and samples (cols)	X	X	X
altExps ^b	experiments with variable number of features		X	X
reducedDims	results of dimensionality reduction		X	X
Rows				
rowData	side information on features	X	X	X
rowPairs	linkage between pairs of associated features		X	X
rowTrees	hierarchical organization of features		X	X
referenceSeq	reference sequences of features		X	X
Columns				
colData	side information on samples	X	X	X
colTrees	hierarchical organization of samples		X	X
sampleMap ^c	flexible sample mapping across experiments			X
Other				
metadata	additional information on the experiment	X	X	X

^aWhile MAE does not contain conventional slots (with the exception of colData), its experiments can include any number of data containers that provide those slots. ^baltExp samples must match assay samples in a 1-to-1 mapping. ^cThe sampleMap supports diverse mappings between samples from different experiments (1-to-1, 1-to-many, many-to-1 and many-to-many). SE, SummarizedExperiment; TreeSE, TreeSummarizedExperiment; MAE, MultiAssayExperiment.

Data resources. Several open microbiome data resources are supported by this framework, enabling direct data import into the aforementioned data containers for further downstream analysis within R/Bioconductor (Table 2). Examples include curatedMetagenomicData³⁷, HoloFood³⁸, microbiomeDataSets (Bioconductor package), and the EBI/MGnify database³⁹. Bioconductor packages include additional built-in demonstration datasets. Collectively, these resources encompass taxonomic

and functional profiles from thousands of published microbiome studies and tens of thousands samples that can be accessed with the available tools for research and educational purposes.

Table 2: Open microbiome data resources (last accessed on 2026-07-05).

Data resource / Package	References	# Studies	# Samples
MGnifyR ^a	39	5,616	429,108
Metalog ^b	40	1,074	157,483
curatedMetagenomicData ^a	37	93	22,588
miaData ^c		9	20,404
microbiomeDataSets ^a		6	19,100
HMP2data ^a		3	11,493
HoloFoodR ^{a,d}	38;41	9	9,990
MicrobiomeBenchmarkData ^a	17	6	1,125

^aBioconductor package with the corresponding name. ^bData retrieval to TreeSE format instructed in OMA online book. ^cZenodo mia community. ^dHoloFood microbiome datasets are deposited in the MGnify database. HMP, Human Microbiome Project.

Alternative frameworks

Several alternative frameworks to microbiome analysis are available. Within Bioconductor, phyloseq²³ provides widely adopted tools with a primary focus on 16S amplicon data analysis. (Tree)SE accommodates all core components of phyloseq objects, including abundance tables, taxonomic annotations, phylogenetic trees, reference sequences, and sample metadata. The advantages of newer (Tree)SE container²⁷ include many additional data elements, a better integration with the overall Bioconductor ecosystem, and enhanced support for multi-assay analyses (see [Data infrastructure](#)). Thus, some originally phyloseq-based tools, such as microbiome R/Bioconductor package¹⁹, have been recently extended with TreeSE support. Related frameworks beyond R/Bioconductor include command-line toolkits supported by active user and developer communities. Mothur²¹ is primarily designed for 16S amplicon analysis and less well suited for integrative metagenomics and multi-omics. QIIME 2²² is a Python-based, AI-ready microbiome multi-omics data science platform. It primarily targets amplicon and metagenomic analyses, and incorporates methods from scikit-bio Python library⁴² which supports multi-omic integration tasks. A particular advantage of these frameworks is their direct connection with the strong machine learning and AI ecosystem in Python, while the interoperability with independently developed omics tools may remain more limited. Anvi'o²⁰ is a complementary, community-driven analysis and visualization platform for microbial 'omics

with a distinct focus on genome-resolved metagenomics and detailed sequence-level analyses. These frameworks emphasize standardized, modular approaches tailored to microbiome analysis and provide validated and user-friendly solutions, supporting efficient end-to-end workflows for microbiome research. Bioconductor emphasizes broad interoperability and shared statistical programming infrastructure with other areas of omics research, and primarily serves the R community. To facilitate migration from alternative frameworks, the OMA book provides conversion tables and utility functions to convert between TreeSE and many external data formats.

Quantitative comparisons indicate enhanced scalability over commonly used alternatives (Figure 3 and Supplementary Figure 1). We executed routine operations on increasingly large subsets of the Metalog database⁴⁰, and measured the execution time and memory consumption with TreeSE, phyloseq, speedyseq, and QIIME 2. Overall, (Tree)SE scaled most efficiently in execution time and memory allocation with increasing feature and sample sizes. It is applicable to abundance matrices of up to 100,000 samples and 10,000 features on a standard laptop, except for the heaviest operations such as phylogenetic beta diversity (UniFrac) calculation, where the general break point is reached around 10,000 samples with all methods. At a feature set size of 10,000 the break point is additionally reached with the PhILR transformation and melting operation. At smaller sample sizes phyloseq or its accelerated extension speedyseq are occasionally faster but these differences are moderate (few seconds on a standard laptop). QIIME 2 is generally slower on lower sample sizes but scales well with increasing sample sizes. Despite the advantages, optimizing time and memory consumption represents a key area for further development.

Data preparation

While several tools for microbiome data analysis are available^{11;12;43}, their implementations typically differ in terms of expected input and output formats. Common data standards promote modular workflows, speeding up methods development, benchmarking, and replication. Building upon these principles, we describe a set of interoperable software libraries that support microbiome data integration and analysis based on the two integrative data containers, TreeSE and MAE (Figure 4).

This framework specifically supports the downstream analysis of high-throughput sequencing data from microbiome experiments, after summarizing the original measurements into abundance tables (*e.g.*, taxonomic or functional profiles). Taxonomic abundance tables are one of the most commonly encountered data types in contemporary microbiome studies. Common measurement platforms include amplicon-based marker gene studies (*e.g.*, 16S rRNA) and metagenomic sequencing, processed with methods such as MetaPhlAn⁴⁴ or Greengenes2⁴⁵. Alternative profiling techniques are available, including phylogenetic microarrays⁴⁶ and high-density optical mapping⁴⁷. Tabular abundance data can also arise from resistome profiling, mapping of other genomic elements, or from functional measurements. The framework is agnostic to the source of abundance data but many of our methodological choices have been guided by data tables that are typical for microbiome profiling studies.




Figure 3: Execution time and allocated memory for routine microbiome data operations. TreeSummarizedExperiment, QIIME 2 and phyloseq—with its extension speedyseq—are three popular frameworks for microbiome analysis. We benchmarked them for five routine operations applied to random subsets of features (rows) and samples (columns) from the Metalog database⁴⁰. Both execution time (s) and allocated memory (MB) were measured as the mean value over ten replicates with varying subsets of the same dimension. Errorbars show the 95% confidence interval of the performance estimate. The benchmarking was conducted in R using the bench library with 4 CPU cores and 16 GB RAM. QIIME 2 lacks allocated memory estimates because memory cannot be monitored for external processes. The source code for the benchmarking analysis is openly available in the Zenodo mia community.

Data import. While the TreeSE data container can be populated with microbiome data from common tabular data types, non-standard file formats may require additional data wrangling. To streamline data import, we provide importers for a variety of standard microbiome file formats, including BIOM⁴⁸, MetaPhlAn/HUMAnN (versions 2-4)⁴⁹, Mothur²¹, QIIME 2²², and the TAXonomic Profile Aggregation and standardization format (taxpasta)⁵⁰. External to Bioconductor, the qiime2R package provides additional functionality to create TreeSE objects directly from QIIME 2 artifacts.

These importers retain provenance of the metadata and the conversion workflow, bridging the gap between the original data processing, other tools, and downstream statistical analysis in R. This allows researchers to focus on interpreting biological patterns rather than wrangling file formats. Moreover, converters between alternative data structures in R are available supporting seamless conversions between TreeSE and BIOM, DADA2⁹, and phyloseq (the `mia::convert*` functions). Imported datasets from different omics modalities can then be interlinked within the MAE data container with its flexible sample mappings as described above. This integrated data object can then be analyzed with dedicated downstream methods.

Similarly, microbiome data stored in TreeSE objects can be exported to external services such as MicrobiomeDB⁵¹ and to other languages (*e.g.*, Julia). These converters facilitate widespread access to microbiome analysis methods.

Quality control and exploration. After importing the data into R, the recommended first step involves basic data exploration and quality control⁵². In addition to methods for identifying contaminant sequences and calculating summary statistics, the framework includes a versatile set of custom visualization methods to assess data variability, outliers, homogeneity of variance and other aspects that may need to be considered in downstream data analysis. The `mia` and `miaViz` packages provide convenient access to many operations for the integrative data containers regarding common tasks in microbiome data exploration and visualization (Figure 5).

Filtering and agglomeration. Microbiome datasets can contain tens of thousands of unique features depending on the sequencing resolution, used reference databases, and the scope of analysis. Moreover, the integration of multiple omics modalities adds another level of challenges. To summarize data into biologically meaningful groups or to reduce multiple comparisons the data is often subsetted, filtered or agglomerated to higher-level taxonomic or functional categories or stratified by sample groups. The package ecosystem, and the `mia` package in particular, feature various tools for such tasks. For instance, the prevalence-based methods can be helpful to remove rare features that are likely sequencing artifacts or present otherwise limited value for the analysis. The mechanism of alternative experiments (`altExp`) allows agglomeration by taxonomic levels or other features such as pathogenicity or functional potential (*e.g.*, butyrate production⁵³, gut-brain modules⁵⁴, and host-associated signatures as in the BugSigDB⁵⁵). The agglomerated data can be incorporated as a new alternative experiment within the original data object. The automated tracking of feature and sample links ensures that changes are propagated across all relevant elements of

the data object, such as transformed abundance tables, phylogenetic trees, or interlinked omics modalities. This helps to avoid redundant processing and enables efficient management of interlinked datasets.

Transformation. Statistical transformations commonly used in microbiome research include relative abundance⁵⁶, phylogenetic ILR (PhILR)⁵⁷, centered log-ratio (CLR), and robust CLR (rCLR)⁵⁸. The latter three aim to mitigate compositionality bias. A related concept that controls for unequal sampling depth is data rarefaction, *i.e.* repeatedly subsampling data to equal read depth and averaging the computed metrics to mitigate information loss⁵⁹ that could arise from subsampling (*rarefying*) data to equal read depth without such iteration^{60–62}. Our framework supports these and a variety of other transformations commonly applied to microbiome data.

Data enrichment. Incorporating additional data, such as transformed feature abundance tables, is a frequent operation in applied data analysis, and it often takes place iteratively. The integrative data containers support gradual and organized accumulation of interlinked data elements including alpha and beta diversity, prevalence calculations, transformations, agglomerated data, reduced dimensions, and other derived data on samples and features. We generally recommend initializing the data objects by pre-calculating such measures already at the beginning of the workflow, where they can then be gradually enriched as the project advances. By systematically extending the upstream data preparation step, instead of ad-hoc calculations scattered amidst the workflow, one can achieve a well-organized and transparent provenance of the integrated data object.




Figure 4: Multi-modal data integration workflow. Following data generation and preprocessing (top two panels), the integrative framework covers five steps of microbiome data integration and analysis: 1) Data is combined into a `TreeSummarizedExperiment` (TreeSE) data container, storing information on the samples (tube icons) and their features (*e.g.*, taxonomic, functional, or resistome profiles; colored circles). TreeSE objects can accommodate multiple assay tables as well as additional hierarchical structures. 2) The TreeSE data object can then be processed to clean the data. This involves filtering out low-quality samples and features, transforming assays for example by normalizing feature-wise to enhance the comparability of sample-specific abundances, or agglomerating features to specific taxonomic levels or functional categories. 3) The cleaned data can then be analyzed using statistical models to uncover relations among samples (*e.g.*, with ordination methods), analyze drivers of variation in community diversity and composition, or test for differential abundance (*e.g.* between sampling time or location). 4) The framework supports the integration of various types of omics modalities through the `MultiAssayExperiment` and omic-specific data containers provided by Bioconductor. For example, the taxonomic and metabolomic profiles of samples can be summarized and compared in a joint analysis among features. 5) Finally, Quarto notebooks aid reproducible analysis and reporting.

Downstream data analysis

Our framework has been designed to improve the management of integrative multi-omic workflows. It facilitates joint operations on interlinked data elements, helping to shift the focus towards multi-modal analyses and methods development while also supporting standard data analyses in microbiome research.

Microbiome data analysis often relies on specialized statistical approaches, since data sparsity, zero-inflation, compositionality, and heteroscedasticity pose challenges for standard methods^{11;56}. Furthermore, microbiome data is hierarchically structured across taxonomic, ecological, temporal, and spatial levels^{7;32;63}. Alternative processing pipelines introduce different biases and generate varying data formats; for example, the same sample can yield differing taxonomic profiles depending on the computational method used⁶⁴. These and other unique characteristics have wide-reaching implications for microbiome analysis, necessitating dedicated analysis methods^{7;11}. The shared data containers support community-driven development and reproducible benchmarking.

The following sections provide an overview of some of the currently available tools, encompassing standard methods for taxonomic, metagenomic, and functional abundance tables as well as integrative analyses between modalities. A comprehensive listing of microbiome analysis tools and resources is maintained in the online book.

Community indicators. Various general indicators are available to quantify the microbial community state or other sample properties. Community diversity is a common measure in microbial ecology¹³ (Figure 5). Alpha diversity metrics generally quantify the observed or estimated number of distinct taxa (*richness*) and

their distribution (*evenness*). Many diversity indices (*e.g.*, Shannon index), combine these two aspects. Phylogenetic diversity metrics additionally account the breadth of phylogenetic relatedness among community members, but are more computationally demanding than abundance-based metrics. A recent implementation, the Stacked Faith’s Phylogenetic Diversity, improves computational efficiency by 1–2 orders of magnitude³⁵ by an optimized single-pass algorithm and utilizing data sparsity. The framework supports analyses with repeated rarefaction through the *mia* package⁵⁹. Other global indices of community state include measures of dominance, rarity, skewness, kurtosis, modality, library size (number of reads), or absolute abundance quantification. Sample information in the TreeSE object (*colData*) can be enriched with these derived quantities. By default, the *mia::addAlpha* method returns a set of complementary alpha diversity measures⁶⁵. Alpha diversity indices have been frequently applied to measure also resistome diversity^{66;67}, and potentially other modalities.

Community similarity. This is commonly evaluated with beta diversity measures, which can be used for quantitative comparisons of multivariate community composition. Our framework supports common dissimilarity metrics in microbial ecology, including Bray-Curtis, Jaccard, Aitchison¹³, and the accelerated phylogeny-aware (striped) UniFrac³⁶. Community dissimilarity is commonly visualized with unsupervised ordination that projects data onto a lower-dimensional space. Among unsupervised techniques, the most commonly used ones include Principal Component Analysis (PCA) and Principal Coordinate Analysis (PCoA) (Figure 5). Statistical tests (*e.g.*, PERMANOVA) and supervised ordination methods, such as distance-based Redundancy Analysis (dbRDA), can complement the unsupervised analyses, and help uncover and quantify covariation between community composition and sample-specific covariates. This type of analysis can benefit from compositionally robust measures (*e.g.*, robust Aitchison distance⁵⁸), or phylogeny-aware dissimilarity measures specifically developed for microbiome research (*e.g.*, UniFrac). Sample dissimilarity provides the basis for many specialized microbiome techniques, including microbial community typing (clustering) with Dirichlet multinomial mixtures⁶⁸ and community state types⁶⁹, or factorization into community signatures⁷⁰. These and other multivariate methods are provided by *mia*, *miaViz*, and additional single-cell packages (*scater*, *scuttle*²⁴). Many omics modalities can benefit from these and other standardized methods.

Association analyses. Differential Abundance Analysis (DAA) and Differential Prevalence Analysis (DPA) methods identify indicator microbial features whose abundances (DAA) or prevalences (DPA) vary across study groups or continuous gradients, for example in terms of spatial or temporal dissimilarity between samples (Figure 5). Many widely used DAA methods support the (Tree)SE framework, including for instance ANCOM-BC2⁷¹, LimROTS²⁶, MaAsLin3⁷², and radEmu⁷³, and benchmarking pipeline⁷⁴. More recently, DPA methods have seen increasing adoption in microbiome research. They ignore abundance data, and focus on comparing feature prevalence *i.e.* presence/absence profiles. Recent implementations include MaAsLin3⁷² and the probabilistic DiPPER⁷⁵. Leveraging the interoperability of the framework, several of these methods have also found applications with other omics

modalities, including transcriptomics⁷⁶, metabolomics⁷², and proteomics²⁶. Standard DAA typically focuses on individual features, ignoring interactions or other covariation among feature abundances. Thus, we support covariation analyses. Recent studies have indicated that elementary methods, such as log-transformed Total Sum Scaled (TSS) counts coupled with a linear model, or presence/absence data with logistic regression, often yield the most consistent results^{16;17}. While agglomerating collinear features as a preprocessing step can reduce multiple testing, increase statistical power, and improve interpretability, it is also prone to subjective parameterization and limited generalizability. Features can be grouped based on clustering results from abundance profiles or on external information drawn from phylogenetic relations, taxonomic classifications or functional properties (*e.g.*, aggregating based on metabolite production⁵³ or disease associations⁵⁵). Beyond these, recent multi-modal causal mediation analysis for microbiome data have been implemented for SE in the multimedia package⁷⁷, highlighting the readiness to support the adoption of the rapidly developing multi-omic methodology.

Function- and sequence-level analyses. Taxonomic analysis is frequently complemented by functional analyses, either based on functional predictions derived from amplicon or metagenomic sequences or from parallel metabolomic, metatranscriptomic or other functional assays. Our framework not only supports the incorporation of functional abundance tables within the same data objects, but many statistical methods are generic and can therefore be applied across different omics data types. Where appropriate, specialized methods are also available; for instance, MaAsLin3⁷² and many other Bioconductor libraries support metatranscriptomics analysis, notame package⁷⁸ for metabolome analyses recently added SE support, and the R for Mass Spectrometry (R4MS) project provides many interoperable tools. The framework has found recent applications in extended analyses of metagenomic features such as antimicrobial resistome composition⁶⁷. We also converted a large-scale compilation of resistome profiles in 8,972 metagenome samples⁶⁶ into the (Tree)SE format to support research and training in this area of microbiome analysis.

Spatio-temporal, ecological, and epidemiological models. Microbiome time series⁷⁹ encompass a wide range of study designs, ranging from short to long follow-up times, sparse to dense data, univariate to multivariate and multi-omic monitoring^{32;80;81}, and real, pseudo, or simulated time series. Moreover, an increasing number of studies incorporate spatial dimension in microbiome analysis⁶³. While spatio-temporal analyses require specialized methods, the miaTime package offers a set of data wrangling methods to support longitudinal, multi-modal microbiome analysis. In addition, Bioconductor contains applicable methods from other research domains, such as single-cell omics and spatial transcriptomics, which use closely related data containers. Spatial and temporal information can be incorporated as a covariate (*e.g.*, location or time) or by estimating temporally or spatially resolved metrics, such as divergence (dissimilarity between consecutive time points), autocorrelation (similarity of closely related samples), and feature modality (whether a feature exhibits multiple abundance states). Another type of temporal analysis includes the time-to-event, or survival models, which are encountered regularly in epidemiological and population cohort studies; for microbiome data, these analyses require specialized approaches, as

proposed in previous studies⁸². We also provide methods to simulate time series from standard models of microbial community dynamics including Hubbell’s neutral model, generalized Lotka-Volterra, and the consumer-resource model through the miaSim package⁸³. Our framework provides a robust basis for developing extended methods for longitudinal multi-omics, which is currently an active area of investigation³². We also converted a large-scale time series of 16,234 gut microbiome profiles from 585 wild baboons over 14 years⁸⁴ into the TreeSE format to support research and training on longitudinal datasets.

Multi-modal integration. A key feature of this framework is to provide the infrastructure supporting the integration of microbial abundance data with other omics modalities, such as transcriptomes, proteomes, or metabolomes^{85;86}, resistomes^{66;67}, host genomes⁸⁷, or environmental exposomes (Figure 5). By leveraging standardized containers (Figure 2), users get fluent access to a growing number of multi-omic methods. Containers reduce the need for manual data wrangling and mitigate the risk of errors (*e.g.*, sample mismatches or data formatting), and support modular and reproducible workflows. They also facilitate incorporation of further downstream tools. Multi-omic methods are particularly relevant for the present ecosystem given its enhanced support for integrative analysis. Examples include association-based approaches, latent variable models, and Machine Learning (ML)-based prediction^{32;88;89}. Our framework is supported by many common multi-omic libraries in Bioconductor: Cross-correlation analysis, provided by mia, is commonly used to capture direct associations of samples or features between two datasets but it does not explicitly account for potential co-variation between microbial features, and methods such as anansi can additionally incorporate known functional relations into association analyses. Unsupervised latent variable methods such as multi-omics factor analysis (MOFA+)⁹⁰ and joint robust PCA (Joint-RPCA)⁹¹ model shared latent structure across modalities, enabling the exploration and visualization of coordinated variation across multiple features and data sources. Supervised methods include for instance IntegratedLearner⁹², which can leverage dependencies between omics layers towards improved prediction performance. These and other methods natively support Bioconductor’s multi-assay framework. Additional tools are provided by multi-omic toolkits such as mixOmics⁹³ provide additional tools, including the discriminant analysis method DIABLO⁹⁴ and other integration approaches. These tools do not yet support the same data containers but can be used seamlessly by using converters or a straightforward extraction of the abundance tables, sample metadata, or other relevant data elements. The unified integrative framework that we have adopted can accelerate methods development, benchmarking, and standardization.

Other methods and programming environments. Several other methods and workflows have been employed in microbiome data science. Our multi-assay framework provides a flexible statistical programming ecosystem dedicated to microbiome data science. It supports the incorporation of more targeted analysis methods and workflows such as SIAMCAT for association testing⁹⁵, mikropml for machine-learning based prediction and classification⁹⁶, NetCoMi for network analysis⁹⁷, miaSim⁸³ for time series simulations, or microbe-set enrichment analyses^{98;99}. These and other similar tools are complementary and they operate primarily on a single feature matrix,

lacking the multi-omic feature selection offered by dedicated integration methods discussed in the previous section. Interoperable containers facilitate the translation of existing tools to adjacent or emerging disciplines.

Our framework can be linked to external microbiome data analysis pipelines using tools such as bioBakery⁴⁹, and foreign platforms through language-specific interfaces, namely Rcpp for C++, reticulate for Python, and JuliaCall for Julia. The overall data analytical paradigm of Bioconductor is similar to that of scikit-bio⁴², a Python-based ecosystem. Whereas scikit-bio benefits from the rapidly evolving machine learning and artificial intelligence tools in Python, Bioconductor’s strengths lie in its vast statistical ecosystem and interoperability with dedicated omics tools and other computational environments. We anticipate that the ongoing developments in cross-language integration will further lower the barriers for cross-utilizing tools across complementary environments. This can be enhanced by additional workflow management tools such as *targets* in R¹⁰⁰ or *Nextflow* for cross-language operations¹⁰¹.

Accessible analysis and community

Modern science is embodied by a critical community capable of assessing discoveries and replicating results¹⁰³. Open source developer communities present a timely manifestation of this principle, with technical, legal, sociological and epistemic consequences³. FAIR principles for software support inclusive access to well-documented and user-friendly methods¹, and are facilitated by the extensive quality control and documentation standards of Bioconductor²⁷.

In addition to these technical demands, the development and quality control of research software rely significantly on community engagement. This stimulates continuous software improvement and promotes the emergence and dissemination of evidence-based best practices^{1;3;79}. Simple approaches should be preferred over complex ones by default, and key analyses should be complemented with visual exploration and verification. Expressive code and meaningful comments support reproducibility and verification. Such practices can greatly improve the transparency and robustness of the analyses and interpretability of the results. OMA facilitates standardization across studies by promoting the use of harmonized data structures, validated methods, and informed default settings. This facilitates methods sharing and benchmarking, helping to establish community best practices and build consensus while allowing users to flexibly select and evaluate methods that are appropriate for their specific data and research questions. The Bioconductor framework naturally aligns with these goals through its open statistical programming ecosystem that encompasses data resources, interoperable software, and educational materials. Community building is actively supported through online training resources and communication channels, making it easier for both developers and users to contribute to the continued critical assessment, development and standardization of the methods and documentation. Together with interactive applications, these resources offer accessible entry points for developers and users regardless of their programming experience.




Figure 5: Microbiome analysis workflows in the Bioconductor ecosystem. The figure demonstrates a range of single-omic and multi-omic analytical approaches available for microbiome research. Using a publicly available multi-omics mouse dataset (taxonomic and metabolomic measurements)⁸⁵, we illustrate: **(a)** taxonomic visualization through phylogenetic trees and abundance heatmaps; **(b–d)** common single-omic analyses including Faith’s phylogenetic diversity³⁵, community composition ordination (dbRDA based on UniFrac dissimilarities¹⁰²), and differential prevalence testing⁷²; and **(e–f)** multi-omic integration strategies combining microbiota and metabolomics data through unsupervised dimensionality reduction (Joint-RPCA⁹¹) to identify shared variation across data layers, and supervised machine learning (IntegratedLearner⁹²) to rank features by their strength of association with experimental conditions or phenotypes of interest.

The OMA book. *Orchestrating Microbiome Analysis with Bioconductor* (OMA) is an online book published at <https://bioconductor.org/books/release/OMA> that provides a comprehensive guide to the (Tree)SE framework for microbiome data science in Bioconductor, including practical guidelines, reproducible workflows, and training material for integrative microbiome analyses. The book is versioned according to the Bioconductor’s biannual release cycle. The work incorporates rich feedback from the user community and microbiome training courses. The OMA book aligns to a broader

Bioconductor convention, where prominent data structures^{27;29;30} designed for multi-table data integration have been described across several domains, such as single-cell omics (OSCA book²⁴) and spatial transcriptomics (OSTA book). These resources use standard data containers to support statistical programming on interlinked multi-table datasets.

Interactive applications. Graphical user interfaces provide accessible entry points. The iSEEtree and miaDash packages provide interactive analysis and visualization tools for microbiome data¹⁰⁴. The former extends iSEE to provide support for TreeSE objects, while the latter builds on iSEEtree to enable analyses without requiring any programming. Their interface supports automatic generation of source code and replicable workflows and analysis templates for non-technical users. Moreover, the tidy R paradigm, available through tidyomics, simplifies the syntax and streamlines analytical workflows¹⁰⁵. These tools allow users to analyze and visualize complex data combinations with ease.

Community. Its widespread adoption has made R one of the most cited software resources of our time². The Bioconductor community is widely recognized for developing open research software for life sciences. Its strengths include minimal documentation standards, portability across operating platforms, quality control (*e.g.* peer-review, unit testing, and continuous integration), usability (APIs), robustness (dependency testing and deprecation mechanisms), and a biannual release cycle with semantic versioning. Over the past two decades Bioconductor has grown into a repository of more than 2,300 actively maintained software packages, including 62 packages related to microbiome research. At the core of the Bioconductor’s data science strategy has been the idea of shared data representations that can enhance the transparency, trustworthiness, and explainability of analyses^{4;31}. This has mitigated redundant efforts and facilitated the translation of methods between research domains^{1;106}. The standards have also played an important role in the formation of broad developer and user communities⁵. Bioconductor has a strong tradition in fostering the community through shared data structures and training activities⁵. An active community of developers and users thrives online, supporting the software sustainability and promoting good practices in scientific computing¹⁰⁷. The community interacts via several online communication channels, including Bioconductor support site, Zulip, GitHub platform, and an email list. The development is influenced by feedback from the user community through these channels, regular meetings, and training workshops.

Outlook

Microbiome research has shifted towards integrative analyses linking various omics towards enhanced mechanistic and causal insights^{13;108}. Technological advances in long-read sequencing, functional assays, spatial profiling, and single-cell analysis have created an increasing demand for integrative techniques that link microbiome data with other omics along varying spatial and temporal dimensions. Future extensions could thus provide guidance for longitudinal and repeated-measures microbiome studies, where temporal dependency and intra-subject correlation introduce additional analytical complexity. The field could also benefit from improved support for the

FAIR data principles¹⁰⁹, for instance by adding built-in refinement, enrichment, and validation methods for community standards^{110–112}.

Moreover, there is a need for integrated multi-omics workflows that can combine microbiome data with other omics to support interpretable hypotheses and causal analyses. In concert with parallel developments in other omics, our framework specifically addresses the need to jointly analyze interlinked abundance tables while supporting hierarchical and tabular side information. Future developments may necessitate the adoption of extended hierarchical data structures such as Bioconductor’s GraphExperiment.

Our framework provides the structured input necessary for probabilistic inference and scalable computation as the recent developments increasingly rely on advanced statistical and ML techniques such as Bayesian modeling and Gaussian processes to accommodate uncertainty and temporal or spatial structures^{7;11}.

Building upon these foundations, the rapid emergence of foundation models and deep learning approaches in biomedicine calls for robust, scalable data representations such as those we have adopted. Applications include transfer learning across studies, interpretable embeddings of microbiome multi-omic profiles, computational prediction of microbial function and metabolites, and integration with AI agents for hypothesis generation and automated analysis, among others¹⁰⁶. The interoperability between R and other programming languages could further facilitate the application of novel ML tools and frameworks.

Open data science infrastructures have become critical components of reproducible research and microbiome applications critically rely on open-source methods^{1;3;79}. To enjoy the benefits of automation while also maintaining flexibility, data scientists often recombine open-source software components into custom workflows. However, the lack of commonly adopted data representation standards can hinder the development, sharing, and application of modular workflows.

Building upon a coordinated community effort, we have reported the development of a comprehensive methods ecosystem for multi-modal microbiome data science in R/Bioconductor, built on widely adopted Bioconductor data containers, software libraries, open data resources, tutorials, and a growing community of developers.

The users can benefit from an extended collection of end-to-end case studies demonstrating how different example analyses can be conducted in practice, from original source data to results. Hence, we have provided comprehensive guidance in the OMA online book, to support independent learning and reproducible analysis in microbiome research and beyond^{5;6}. We welcome new community contributions of well-documented methods and workflow demonstrations. As the field evolves, the developer community will continue to integrate new tools and approaches to improve integrative analyses of microbiome data.

Data availability

All datasets used in this article are publicly available with a permanent digital object identifier and an open source license from the specified sources.

Code availability

This article presented v1.1.1 of the OMA book (2026-07-05) published at <https://bioconductor.org/books/release/OMA>, which depends on R v4.6 and Bioconductor v3.24. Analysis scripts and results of the benchmarking analysis are available on Zenodo (mia community).

References

- [1] Barker, M. *et al.* Introducing the FAIR principles for research software. *Scientific Data* **9**, 622 (2022).
- [2] Pearson, H., Ledford, H., Hutson, M. & Van Noorden, R. The most-cited papers of the twenty-first century. *Nature* **640**, 588–593 (2025).
- [3] Hocquet, A. *et al.* Software in science is ubiquitous yet overlooked. *Nature Computational Science* **4**, 465–468 (2024).
- [4] Gentleman, R. *et al.* Bioconductor: Open software development for computational biology and bioinformatics. *Genome Biology* **5**, R80 (2004).
- [5] Drnevich, J. *et al.* Learning and teaching biological data science in the Bioconductor community. *PLoS Computational Biology* **21**, e1012925 (2025).
- [6] Timmis, K. *et al.* A concept for international societally relevant microbiology education and microbiology knowledge promulgation in society. *Microbial Biotechnology* **17**, e14456 (2024).
- [7] Moreno-Indias, I. *et al.* Statistical and machine learning techniques in human microbiome studies: Contemporary challenges and solutions. *Frontiers in Microbiology* **12**, 635781 (2021).
- [8] Ragan-Kelley, B. *et al.* Collaborative cloud-enabled tools allow rapid, reproducible biological insights. *The ISME Journal* **7**, 461–464 (2013).
- [9] Callahan, B. J. *et al.* DADA2: High-resolution sample inference from illumina amplicon data. *Nature Methods* **13**, 581–583 (2016).
- [10] Holmes, S. Successful strategies for human microbiome data generation, storage and analyses. *Journal of Biosciences* **44**, 111 (2019).
- [11] Jeganathan, P. & Holmes, S. P. A statistical perspective on the challenges in molecular microbial biology. *Journal of Agricultural, Biological and Environmental Statistics* **26**, 131–160 (2021).
- [12] Willis, A. D. Rigorous statistical methods for rigorous microbiome science. *mSystems* **4**, e00117–19 (2019).

- [13] Bastiaanssen, T. F. S., Quinn, T. P. & Loughman, A. Bugs as features (part 1): concepts and foundations for the compositional data analysis of the microbiome–gut–brain axis. *Nature Mental Health* **1**, 930–938 (2023).
- [14] Bastiaanssen, T. F. S., Quinn, T. P. & Loughman, A. Bugs as features (part 2): a perspective on enriching microbiome–gut–brain axis analyses. *Nature Mental Health* **1**, 939–949 (2023).
- [15] Nearing, J. T. *et al.* Microbiome differential abundance methods produce different results across 38 datasets. *Nature Communications* **13**, 342 (2022).
- [16] Peltó, J., Auranen, K., Kujala, J. & Lahti, L. Elementary methods provide more replicable results in microbial differential abundance analysis. *Briefings in Bioinformatics* **26**, bbaf130 (2025).
- [17] Gamboa-Tuz, S. D. *et al.* Commonly used compositional data analysis implementations are not advantageous in microbial differential abundance analyses benchmarked against biological ground truth. *bioRxiv* (2025).
- [18] Wen, T. *et al.* The best practice for microbiome analysis using R. *Protein & Cell* **14**, 713–725 (2023).
- [19] Shetty, S. A. & Lahti, L. Microbiome data science. *Journal of Biosciences* **44**, 115 (2019).
- [20] Eren, A. M. *et al.* Community-led, integrated, reproducible multi-omics with Anvi'o. *Nature Microbiology* **6**, 3–6 (2021).
- [21] Schloss, P. D. *et al.* Introducing mothur: Open-source, platform-independent, community-supported software for describing and comparing microbial communities. *Applied and Environmental Microbiology* **75**, 7537–7541 (2009).
- [22] Bolyen, E. *et al.* Reproducible, interactive, scalable and extensible microbiome data science using QIIME 2. *Nature Biotechnology* **37**, 852–857 (2019).
- [23] McMurdie, P. J. & Holmes, S. phyloseq: a Bioconductor package for handling and analysis of high-throughput phylogenetic sequence data. *Pacific Symposium on Biocomputing* 235–246 (2012).
- [24] Amezquita, R. A. *et al.* Orchestrating single-cell analysis with Bioconductor. *Nature Methods* **17**, 137–145 (2020).
- [25] Crowell, H. L. *et al.* Orchestrating spatial transcriptomics analysis with bioconductor. *bioRxiv* (2026).
- [26] Anwar, A. M., Jeba, A., Lahti, L. & Coffey, E. LimROTS: a hybrid method integrating empirical Bayes and reproducibility-optimized statistics for robust differential expression analysis. *Bioinformatics* **41**, btaf570 (2025).

- [27] Huber, W. *et al.* Orchestrating high-throughput genomic analysis with Bioconductor. *Nature Methods* **12**, 115–121 (2015).
- [28] Righelli, D. *et al.* SpatialExperiment: infrastructure for spatially-resolved transcriptomics data in R using Bioconductor. *Bioinformatics* **38**, 3128–3131 (2022).
- [29] Huang, R. *et al.* TreeSummarizedExperiment: a S4 class for data with hierarchical structure. *F1000Res* **9**, 1246 (2021).
- [30] Ramos, M. *et al.* Software for the integration of multiomics experiments in Bioconductor. *Cancer research* **77**, e39–e42 (2017).
- [31] Lawrence, M. *et al.* Software for computing and annotating genomic ranges. *PLoS Computational Biology* **9**, e1003118 (2013).
- [32] Sherwani, M. K. *et al.* Multi-omics time-series analysis in microbiome research: a systematic review. *Briefings in Bioinformatics* **26**, bba502 (2025).
- [33] Husso, A. *et al.* Impacts of maternal microbiota and microbial metabolites on fetal intestine, brain, and placenta. *BMC Biology* **21**, 207 (2023).
- [34] Lun, A. T. L., Pagès, H. & Smith, M. L. beachmat: A Bioconductor C++ API for accessing high-throughput biological data from a variety of R matrix types. *PLoS Computational Biology* **14**, e1006135 (2018).
- [35] Armstrong, G. *et al.* Efficient computation of Faith’s phylogenetic diversity with applications in characterizing microbiomes. *Genome Research* **31**, 2131–2137 (2021).
- [36] McDonald, D. *et al.* Striped UniFrac: enabling microbiome analysis at unprecedented scale. *Nature Methods* **15**, 847–848 (2018).
- [37] Pasolli, E. *et al.* Accessible, curated metagenomic data through ExperimentHub. *Nature Methods* **14**, 1023–1024 (2017).
- [38] Rogers, A. B. *et al.* HoloFood data portal: holo-omic datasets for analysing host–microbiota interactions in animal production. *Database (Oxford)* **2025**, baae112 (2025).
- [39] Gurbich, T. A. *et al.* MGnify genomes: A resource for biome-specific microbial genome catalogues. *Journal of molecular biology* **435**, 168016 (2023).
- [40] Kuhn, M. *et al.* Metalog: curated and harmonised contextual data for global metagenomics samples. *Nucleic Acids Research* **54**, D826–D834 (2026).
- [41] Borman, T. *et al.* HoloFoodR: a statistical programming framework for holo-omics data integration workflows. *Bioinformatics* **41**, btaf605 (2025).

- [42] Aton, M. *et al.* Scikit-bio: a fundamental python library for biological omic data analysis. *Nature Methods* **23**, 274–276 (2026).
- [43] Marcos-Zambrano, L. J. *et al.* A toolbox of machine learning software to support microbiome analysis. *Frontiers in Microbiology* **14** (2023).
- [44] Blanco-Míguez, A. *et al.* Extending and improving metagenomic taxonomic profiling with uncharacterized species using MetaPhlAn 4. *Nature Biotechnology* **41**, 1633–1644 (2023).
- [45] McDonald, D. *et al.* Greengenes2 unifies microbial data in a single reference tree. *Nature Biotechnology* **42**, 715–718 (2024).
- [46] Rajilić-Stojanović, M. *et al.* Development and application of the human intestinal tract chip, a phylogenetic microarray: analysis of universally conserved phylotypes in the abundant microbiota of young and elderly adults. *Environmental Microbiology* **11**, 1736–1751 (2009).
- [47] Abakumov, S. *et al.* DynaMAP: Dynamic microbiome abundance profiling through high density optical mapping. *bioRxiv* (2024).
- [48] McDonald, D. *et al.* The biological observation matrix (BIOM) format or: how i learned to stop worrying and love the ome-ome. *Gigascience* **1**, 7 (2012).
- [49] Beghini, F. *et al.* Integrating taxonomic, functional, and strain-level profiling of diverse microbial communities with bioBakery 3. *elife* **10**, e65088 (2021).
- [50] Beber, M. E., Borry, M., Stamouli, S. & Yates, J. A. F. TAXPASTA: TAXonomic profile aggregation and STAndardisation. *Journal of Open Source Software* **8**, 5627 (2023).
- [51] Oliveira, F. S. *et al.* MicrobiomeDB: a systems biology platform for integrating, mining and analyzing microbiome experiments. *Nucleic Acids Research* **46**, D684–D691 (2018).
- [52] Zuur, A. F., Ieno, E. N. & Elphick, C. S. A protocol for data exploration to avoid common statistical problems. *Methods in Ecology and Evolution* **1**, 3–14 (2010).
- [53] Kullberg, R. F. J. *et al.* Association between butyrate-producing gut bacteria and the risk of infectious disease hospitalisation: results from two observational, population-based microbiome studies. *The Lancet Microbe* **5**, 100864 (2024).
- [54] Valles-Colomer, M. *et al.* The neuroactive potential of the human gut microbiota in quality of life and depression. *Nature Microbiology* **4**, 623–632 (2019).

- [55] Geistlinger, L. *et al.* BugSigDB captures patterns of differential abundance across a broad range of host-associated microbial signatures. *Nature Biotechnology* **42**, 790–802 (2024).
- [56] Gloor, G. B., Macklaim, J. M., Pawlowsky-Glahn, V. & Egozcue, J. J. Microbiome datasets are compositional: And this is not optional. *Frontiers in Microbiology* **8** (2017).
- [57] Silverman, J. D., Washburne, A. D., Mukherjee, S. & David, L. A. A phylogenetic transform enhances analysis of compositional microbiota data. *eLife* **6**, e21887 (2017).
- [58] Martino, C. *et al.* A novel sparse compositional technique reveals microbial perturbations. *mSystems* **4**, e00016–19 (2019).
- [59] Schloss, P. D. Rarefaction is currently the best approach to control for uneven sequencing effort in amplicon sequence analyses. *mSphere* **9**, e00354–23 (2024).
- [60] McMurdie, P. J. & Holmes, S. Waste Not, Want Not: Why rarefying microbiome data is inadmissible. *PLoS Computational Biology* **10**, e1003531 (2014).
- [61] Schloss, P. D. Waste not, want not: revisiting the analysis that called into question the practice of rarefaction. *mSphere* **9**, e00355–23 (2024).
- [62] Willis, A. D. Rarefaction, alpha diversity, and statistics. *Frontiers in Microbiology* **10**, 2407 (2019).
- [63] Ruuskanen, M. O., Sommeria-Klein, G., Havulinna, A. S., Niiranen, T. J. & Lahti, L. Modelling spatial patterns in host-associated microbial communities. *Environmental Microbiology* **23**, 2374–2388 (2021).
- [64] Quince, C., Walker, A. W., Simpson, J. T., Loman, N. J. & Segata, N. Shotgun metagenomics, from sampling to analysis. *Nature Biotechnology* **35**, 833–844 (2017).
- [65] Cassol, I., nez, M. I. & Bustamante, J. P. Key features and guidelines for the application of microbial alpha diversity metrics. *Scientific Reports* **15**, 622 (2025).
- [66] Lee, K. *et al.* Population-level impacts of antibiotic usage on the human gut microbiome. *Nature Communications* **14**, 1191 (2023).
- [67] Salehi, M. *et al.* Gender differences in global antimicrobial resistance. *npj Biofilms and Microbiomes* **11**, 79 (2025).
- [68] Holmes, I., Harris, K. & Quince, C. Dirichlet multinomial mixtures: Generative models for microbial metagenomics. *PLoS One* **7**, e30126 (2012).

- [69] DiGiulio, D. B. *et al.* Temporal and spatial variation of the human microbiota during pregnancy. *Proceedings of the National Academy of Sciences* **112**, 11060–11065 (2015).
- [70] Frioux, C. *et al.* Enterosignatures define common bacterial guilds in the human gut microbiome. *Cell Host & Microbe* **31**, 1111–1125.e6 (2023).
- [71] Lin, H. & Peddada, S. D. Multigroup analysis of compositions of microbiomes with covariate adjustments and repeated measures. *Nature Methods* **21**, 83–91 (2024).
- [72] Nickols, W. A. *et al.* MaAsLin 3: Refining and extending generalized multi-variable linear models for meta-omic association discovery. *Nature Methods* **23**, 554–564 (2026).
- [73] Clausen, D. S., Teichman, S. V. & Willis, A. D. Estimating ratios of means of multicategory data observed with sample and category perturbations. *Biometrika* (2026).
- [74] Calgaro, M., Romualdi, C., Risso, D. & Vitulo, N. benchdamic: benchmarking of differential abundance methods for microbiome data. *Bioinformatics* **39**, btac778 (2022).
- [75] Pelto, J., Auranen, K., Kujala, J. V. & Lahti, L. A Bayesian approach to differential prevalence analysis with applications in microbiome studies (2026). [2602.05938](#).
- [76] Love, M. I., Huber, W. & Anders, S. Moderated estimation of fold change and dispersion for rna-seq data with deseq2. *Genome Biology* **15**, 550 (2014).
- [77] Jiang, H. *et al.* Multimedia: multimodal mediation analysis of microbiome data. *Microbiology Spectrum* **13**, e01131–24 (2025).
- [78] Klåvus, A. *et al.* “Notame”: Workflow for non-targeted LC–MS metabolic profiling. *Metabolites* **10** (2020).
- [79] Berg, G. *et al.* Microbiome definition re-visited: old concepts and new challenges. *Microbiome* **8**, 103 (2020).
- [80] Velten, B. *et al.* Identifying temporal and spatial patterns of variation from multimodal data using MEFISTO. *Nature Methods* **19**, 179–186 (2022).
- [81] Fukuyama, J. *et al.* Multidomain analyses of a longitudinal human microbiome intestinal cleanout perturbation experiment. *PLoS Computational Biology* **13**, e1005706 (2017).
- [82] Salosensaari, A. *et al.* Taxonomic signatures of cause-specific mortality risk in human gut microbiome. *Nature Communications* **12**, 2671 (2021).

- [83] Gao, Y. *et al.* miaSim: an R/Bioconductor package to easily simulate microbial community dynamics. *Methods in Ecology and Evolution* **14**, 1967–1980 (2023).
- [84] Grieneisen, L. *et al.* Gut microbiome heritability is nearly universal but environmentally contingent. *Science* **373**, 181–186 (2021).
- [85] Hintikka, J. *et al.* Xylo-oligosaccharides in prevention of hepatic steatosis and adipose tissue inflammation: Associating taxonomic and metabolomic patterns in fecal microbiomes with biclustering. *International Journal of Environmental Research and Public Health* **18** (2021).
- [86] Mangnier, L. *et al.* A systematic benchmark of integrative strategies for microbiome-metabolome data. *Communications Biology* **8** (2025).
- [87] Qin, Y. *et al.* Combined effects of host genetics and diet on human gut microbiota and incident disease in a single population cohort. *Nature Genetics* **54**, 134–142 (2022).
- [88] Sankaran, K. & Holmes, S. P. Multitable methods for microbiome data integration. *Frontiers in Genetics* **10**, 627 (2019).
- [89] Sankaran, K. & Holmes, S. P. Latent variable modeling for the microbiome. *Biostatistics* **20**, 599–614 (2019).
- [90] Argelaguet, R. *et al.* MOFA+: a statistical framework for comprehensive integration of multi-modal single-cell data. *Genome Biology* **21**, 111 (2020).
- [91] Vargas, B. C. *et al.* Joint-RPCA: Domain-aware multi-omics integration for systems microbiology. *Molecular Systems Biology* (2026). In press.
- [92] Mallick, H. *et al.* An integrated Bayesian framework for multi-omics prediction and classification. *Statistics in Medicine* **43**, 983–1002 (2024).
- [93] Rohart, F., Gautier, B., Singh, A. & Lê Cao, K.-A. mixOmics: An R package for ‘omics feature selection and multiple data integration. *PLoS Computational Biology* **13**, e1005752 (2017).
- [94] Singh, A. *et al.* Diablo: an integrative approach for identifying key molecular drivers from multi-omics assays. *Bioinformatics* **35**, 3055–3062 (2019).
- [95] Wirbel, J. *et al.* Microbiome meta-analysis and cross-disease comparison enabled by the SIAMCAT machine learning toolbox. *Genome Biology* **22**, 93 (2021).
- [96] Topçuoğlu, B. D. *et al.* mikropml: User-friendly R package for supervised machine learning pipelines. *Journal of Open Source Software* **6**, 3073 (2021).
- [97] Peschel, S., Müller, C. L., von Mutius, E., Boulesteix, A.-L. & Depner, M. Net-CoMi: network construction and comparison for microbiome data in R. *Briefings*

- in *Bioinformatics* **22**, bbaa290 (2021).
- [98] Kou, Y., Xu, X., Zhu, Z., Dai, L. & Tan, Y. Microbe-set enrichment analysis facilitates functional interpretation of microbiome profiling data. *Scientific Reports* **10**, 21466 (2020).
 - [99] Pham, C. M., Rankin, T. J., Stinear, T. P., Walsh, C. J. & Ryan, F. J. TaxSEA: rapid interpretation of microbiome alterations using taxon set enrichment analysis and public databases. *Briefings in Bioinformatics* **26**, bbaf173 (2025).
 - [100] Landau, W. M. The targets r package: a dynamic make-like function-oriented pipeline toolkit for reproducibility and high-performance computing. *Journal of Open Source Software* **6**, 2959 (2021). URL <https://doi.org/10.21105/joss.02959>.
 - [101] Di Tommaso, P. *et al.* Nextflow enables reproducible computational workflows. *Nature Biotechnology* **35**, 316–319 (2017).
 - [102] Lozupone, C. & Knight, R. UniFrac: a new phylogenetic method for comparing microbial communities. *Applied Environmental Microbiology* **71**, 8228–8235 (2005).
 - [103] Wootton, D. *The Invention of Science: A New History of the Scientific Revolution* First u.s. edition edn (Harper, an imprint of HarperCollins Publishers, New York, 2015). Paperback edition published December 13, 2016 by Harper Perennial.
 - [104] Benedetti, G. *et al.* iSEETree: interactive explorer for hierarchical data. *Bioinformatics Advances* **5**, vbaf107 (2025).
 - [105] Hutchison, W. J. *et al.* The tidyomics ecosystem: enhancing omic data analyses. *Nature Methods* **21**, 1166–1170 (2024).
 - [106] D’Elia, D. *et al.* Advancing microbiome research with machine learning: key findings from the ml4microbiome COST action. *Frontiers in Microbiology* **14**, 1257002 (2023).
 - [107] Wilson, G. *et al.* Good enough practices in scientific computing. *PLoS Computational Biology* **13**, e1005510 (2017).
 - [108] Joos, R. *et al.* Examining the healthy human microbiome concept. *Nature Reviews Microbiology* **23**, 192–205 (2024).
 - [109] Wilkinson, M. D. *et al.* The FAIR guiding principles for scientific data management and stewardship. *Scientific Data* **3**, 160018 (2016).
 - [110] Field, D. *et al.* The minimum information about a genome sequence (migs) specification. *Nature Biotechnology* **26**, 541–547 (2008).

- [111] Mirzayi, C. *et al.* Reporting guidelines for human microbiome research: the STORMS checklist. *Nature Medicine* **27**, 1885–1892 (2021).
- [112] Kelliher, J. M., Mirzayi, C., Bordenstein, S. R. *et al.* STREAMS guidelines: standards for technical reporting in environmental and host-associated microbiome studies. *Nature Microbiology* **10**, 3059–3068 (2025).

Consortia

OMA consortium: Arora, Nalin; Bastiaanssen, Thomaz; Beber, Moritz E.; Bektanov, Aituar; Benedetti, Giulio; Bhowmick, Sabuj; Borman, Tuomas; Braccia, Domenick J.; Calgaro, Matteo; Callan, Danielle; Chaudhary, Jiya; Cordazzo Vargas, Bianca; Corrada Bravo, Héctor; Courbayre Dussau, Basil; Das, Sneha; Eckermann, Henrik; Erawijantari, Pande; Ernst, Felix G.M.; Gaudron-Parry, Elliot; Gopalakrishnan, Shyam; Hakula, Maija; Hanhineva, Kati; Havulinna, Aki; Hindström, Rasmus; Jeba, Akewak; Lahti, Leo; Limborg, Morten; Lindgren, Himmi; Liu, Yihan; Mallick, Himel; Martino, Cameron; Mongad, Dattatray; Müller, Christian L.; Mulu, Geraldson; Shen-hay, Liat; Nynäs, Signe; Pasanen, Jesse; Paulo Casucci dos Santos, Joao; Pelto, Juho; Peschel, Stefanie; Potbhare, Renuka; Pralas, Théotime; Ramos, Marcel; Raulo, Aura; Sannikov, Artur; Seraidarian, Ely; Serizay, Jacques; Shetty, Sudarshan; Shigdel, Rajesh; Suksi, Vilhelm; Tammi, Eetu; Valderrama, Benjamin; Waldron, Levi; Yang, Lu.

Author contributions

L.L. and T.Bo. conceived and coordinated the study. L.L., T.Bo., G.B., A.R., A.S., R.H. and T.Ba. performed the data analysis. L.L., T.Bo. and G.B. wrote the original draft. L.L. acquired funding. All authors contributed material, participated in manuscript preparation and approved the final version.

Competing interests

The authors declare no competing interests.

Acknowledgments

We thank CSC – IT Center for Science, Finland, for computational resources. This project received funding from the European Union’s Horizon 2020 research and innovation programme under grant agreement No 952914 and Research Council of Finland (decision 374944). TB was additionally supported by the University of Turku Graduate School. We thank Christopher Quince and Kihyun Lee for providing resistome profiling data, Justin Silverman for PhILR support, Ben Allen and Sandy Rogers for MGnifyR support, Elizabeth Archie and Johannes Björk for support with the wild baboon data set, and the following individuals for supporting the project: Richa Ashma; Nitin Bayal; Chouaib Benchakra; Karoline Faust; Robert D. Finn; Samuel Hillman; Ruizhu Huang; Lorna Richardson; Daniel Rios Garza; Matti Ruuskanen; Mahkameh Salehi. We extend our thanks to all contributors of the open-source resources used in this work.